

HackSprint

24-Hour Hackathon

17-18
OCTOBER

VENUE DETAILS
Manipal Academy of Higher
Education (MAHE)

SATURDAY
09 A.M. Onwards

YOUR NAME: Pratham Pai

TEAM NAME: recurSON

TRACK NO: T.03

**REGISTER
NOW**

Team Size
2 Members

ORGANIZED BY:
commodle

Solution

THE PROBLEM

You approved the tool. You didn't approve what the model reads.

AI agents load tools from third-party MCP servers. The model reads every tool's full description and schema, but the human approves only a short rendered summary. Whatever sits in that gap is an attack surface.

- ! **Tool poisoning** Hidden instructions in a tool description: “read ~/.ssh and send it here”.
- ? **Invisible text** Unicode tag characters render as nothing on screen but reach the model in full.
- ↻ **Rug pull** The tool silently changes its definition after you approved it.

WHY NOW


Apr 2025: Invariant Labs showed a poisoned MCP tool making Cursor leak SSH keys, and that tools can change after approval.

OWASP MCP Top 10 lists Tool Poisoning as MCP03:2025.

OUR SOLUTION · FILIGREE *Every last character, accounted for.*

Approval Integrity for AI-agent tools

A tool may run only while all three representations are the same object, checked at every call:



Any mismatch → **APPROVAL INVALID** → **BLOCK**

- 1 **Capture** the exact tool object handed to the agent, with nothing stripped or normalised.
- 2 **Reveal** it side by side with what the human sees, with hidden characters decoded.
- 3 **Analyze** with a deterministic scanner first, then an advisory LLM.
- 4 **Bind** approval to a SHA-256 fingerprint and re-check it before every call.

Tech Stack & Architecture

STACK

Security core
Python 3.11 · FastAPI · Pydantic v2 · hashlib SHA-256

MCP
Official MCP Python SDK over Streamable HTTP (minimal JSON-RPC fallback)


AI analysis
Provider-agnostic LLM adapter with structured outputs · ReplayAnalyzer for offline demos

Storage
SQLite behind a small interface (Postgres-ready) · SQLAlchemy


Frontend
React + TypeScript + Vite · Tailwind · Lucide icons

Testing
pytest · React Testing Library · Playwright golden path

ARCHITECTURE



Security Pipeline



DECISION POLICY priority runs top to bottom


Capture or integrity failure	BLOCK
No approval exists	BLOCK · NO_APPROVAL
Fingerprint ≠ approved fingerprint	BLOCK · MANIFEST_DRIFT
High / critical finding	BLOCK · no override
Medium finding	REVIEW + confirmation
Analyzer unavailable	REVIEW · scanner still enforced
Clean and integrity valid	Eligible for human approval

THREE INVARIANTS

- Integrity invariant**
A call runs only if the freshly captured manifest equals the approved one and policy still permits it.
- Representation invariant**
The approval UI is generated from the same canonical object the agent's registry receives.
- Advisor invariant**
The analyzer can be wrong or manipulated; integrity never depends on it being right.

Data Flow Diagram

① APPROVAL TIME



② CALL TIME (every tools/call)




TOOL TRUST STATES



Demo of our Project

A The gap: what you see vs what the model receives




Prototype UI mockups built from our design spec. All secrets and endpoints are mock data. The product might change later.

Identical on screen


98 hidden chars decoded

Blocked by scanner rules, not LLM opinion

B The block: attack fails before execution



C The rug pull: approval revoked on drift



Others

1

Gateways already filter, fingerprint and intercept. We don't claim those primitives.

Our angle: make the approved object itself explicit and verifiable.

- exact model-visible representation
- human view derived from that same object
- deterministic + semantic analysis of it
- cryptographic approval binding
- call-time revalidation

Category: Agent Approval Integrity. It complements a gateway rather than replacing one.

2

Three sets, frozen before tuning:

- Authored: ~10 poisoned, 15 benign, 5 rug-pull
- Public held-out: published attack samples, each sourced
- Evasive: homoglyphs, invisible chars, multilingual, split across fields

Reported for scanner-only vs LLM-only vs combined:

recall, benign acceptance, false-positive rate, p50/p95 latency.

We show real false negatives on stage instead of hiding them.

3

Out of scope today:

- malicious server behaviour with an unchanged definition
- instructions arriving in tool results
- clients that bypass the instrumented path
- subtle analyzer misses (scanner + human review mitigate them)

Next:

signed manifests & server provenance · proxy mode for any MCP client · approval expiry · per-request pinning · response-level integrity

“Filigree doesn't ask whether a tool is safe. It verifies the AI can only act on what the human actually approved. WYSIWYS.”

<https://github.com/prathampai-college/filigree>